

# Full Stack Development

## Containers, Microservices and UI

## 2. Docker Containers

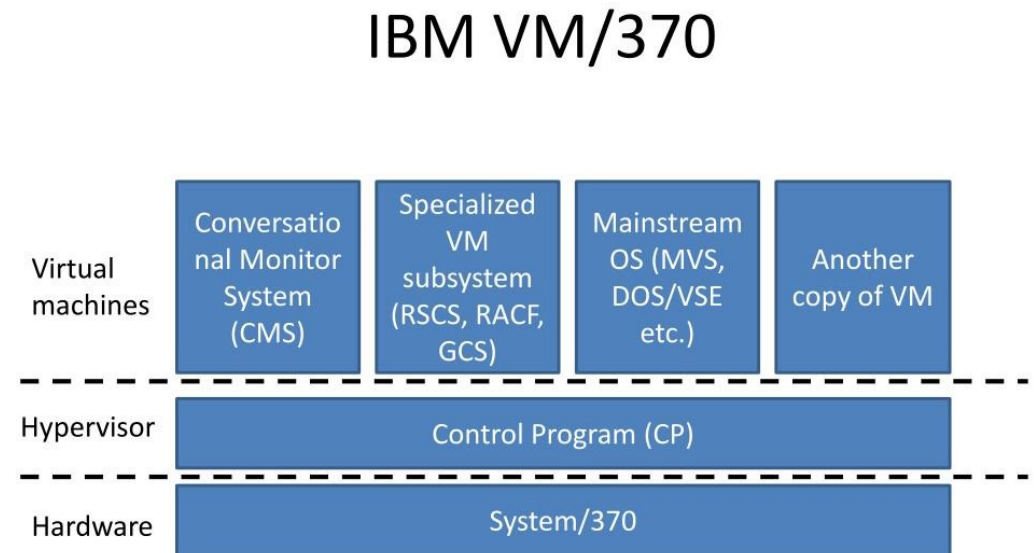


# Virtualization

- The massive increase in hardware computing power resulted software not being able to use all of the available resources
- In order to optimize use of hardware, several virtual machines could be run on the same hardware
- This virtualization model enabled the development of cloud computing
- Developers are completely insulated from the hardware
  - Code is written to run in a virtual environment
  - Virtual environments can be described in a declarative language like terraform or cloud formation in AWS for example

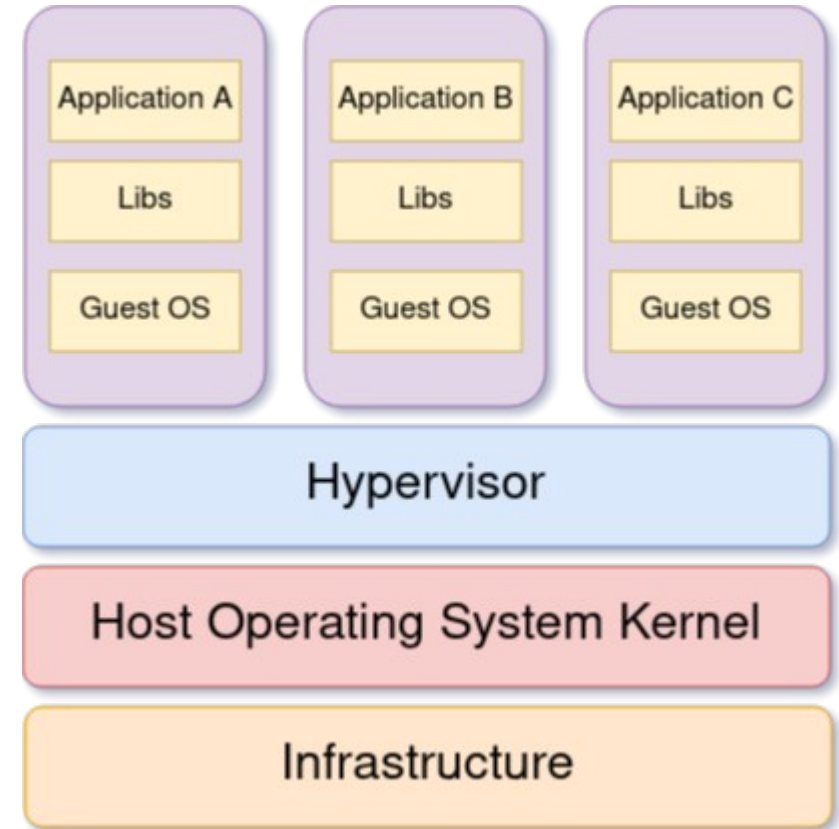
# Virtualization Origins

- In the 1970s, IBM wanted to replace multiple hardware lines and associated operating systems with the 370
- This would have required clients to do massive rewrites of their code
- The VM operating system used a hypervisor to emulate the legacy hardware and software in a VM running on a 370
- This allowed for a smooth transition for clients from legacy systems to the 370




# Virtual Machines

- Each VM is a full installation of a complete OS
  - The VM hard drive maps to a file or files on the host OS
  - The guest OS hardware calls are relayed to the host OS by the hypervisor
- VMs are slow to start and have a large footprint
  - Great for emulating a computer
  - Too heavyweight for running a small lightweight process




# Containers

- A container is a lightweight process managed by the Docker engine
  - It has no persistent storage
  - Contains only what is needed to run the application
  - Small footprint
  - Fast start up and shutdown
- Based on Linux containers
  - Use specific features of the Linux kernel
  - Windows can run containers by using an embedded Linux VM
    - WLS - “Windows Linux subsystem”

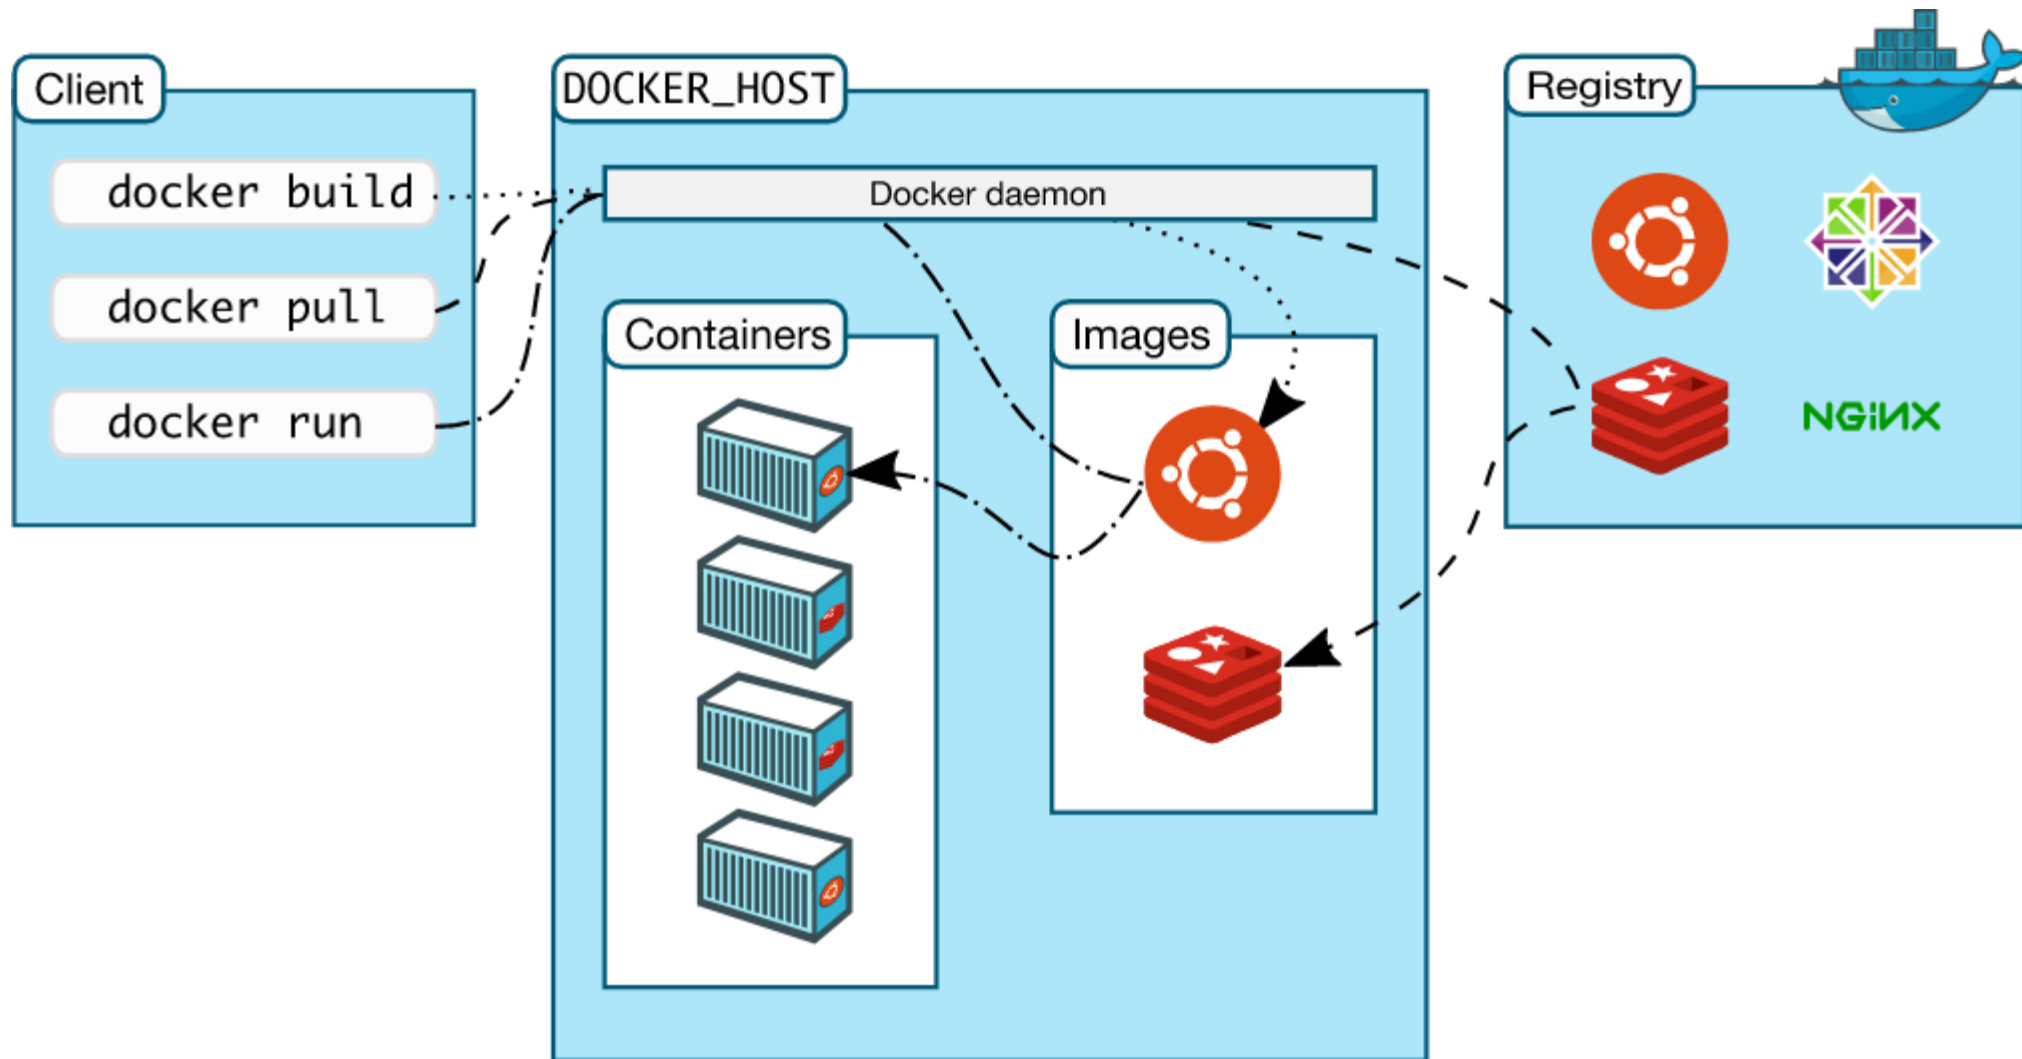


# Implementation

- Linux containers run in isolated environments using Linux namespaces and control groups
- Provide resource limitation, prioritization, accounting, and control
- Hides the process space and resource information of each container from the others
- Docker is an implementation of Linux containers



# Docker Architecture



# Docker Terminology

- Docker Daemon or Engine
  - The process that manages images and containers on the Docker host
  - The Docker CLI is used to request services from the Docker engine
- Docker Image
  - Analogous to an executable file - template for running a container
- Docker Container
  - A Docker image that is executing or has finished executing
  - Analogous to a process that is running an executable file
  - Multiple containers can be created and run from a single image
- Docker Registry
  - A version collection Docker images
  - Each set of versions for an image is referred to as an image repository



# Docker Images

- Docker images are read only
  - Uniquely identified by hash codes
- Built-up in layers
  - Uses Linux union file system, also referred to as an overlay file system
  - Each layer is immutable identified by a unique hash code
  - Layers are shared by images – only one copy of a layer exists



Docker container  
(AUFS storage-driver demonstrating whiteout file)

# Docker Registry

- The local registry cache
  - On the machine running Docker
  - Images pulled from other registries are cached here
  - This is the first registry searched for a requested image
- Docker Hub
  - Public repository maintained by Docker
  - Searched by default after the local registry
- Other registries
  - Docker can be configured to use other registries
  - Allows control over which images Docker pulls
  - Ensures only approved images are used by Docker installation

# Docker Repository

- Images are versioned
  - A set of versioned images is called repository
  - A specific image is referenced by <image\_name>:<version\_tag>
  - The following are different versions of the Ubuntu image
    - ubuntu:18.04
    - ubuntu:20.04
  - If no version tag is specified, then the version defaults to “latest”
    - Pulling the image **ubuntu** is the same as pulling **ubuntu:latest**
- Images are uniquely identified by their digest value
  - Tags are identifies that are added for convenience
  - Images do not have to have tags but a single image can have multiple tags
  - Images can only be deleted if they have zero or one tags



# Exploring Docker Registries

Demo





# Working with Docker Images

Demo





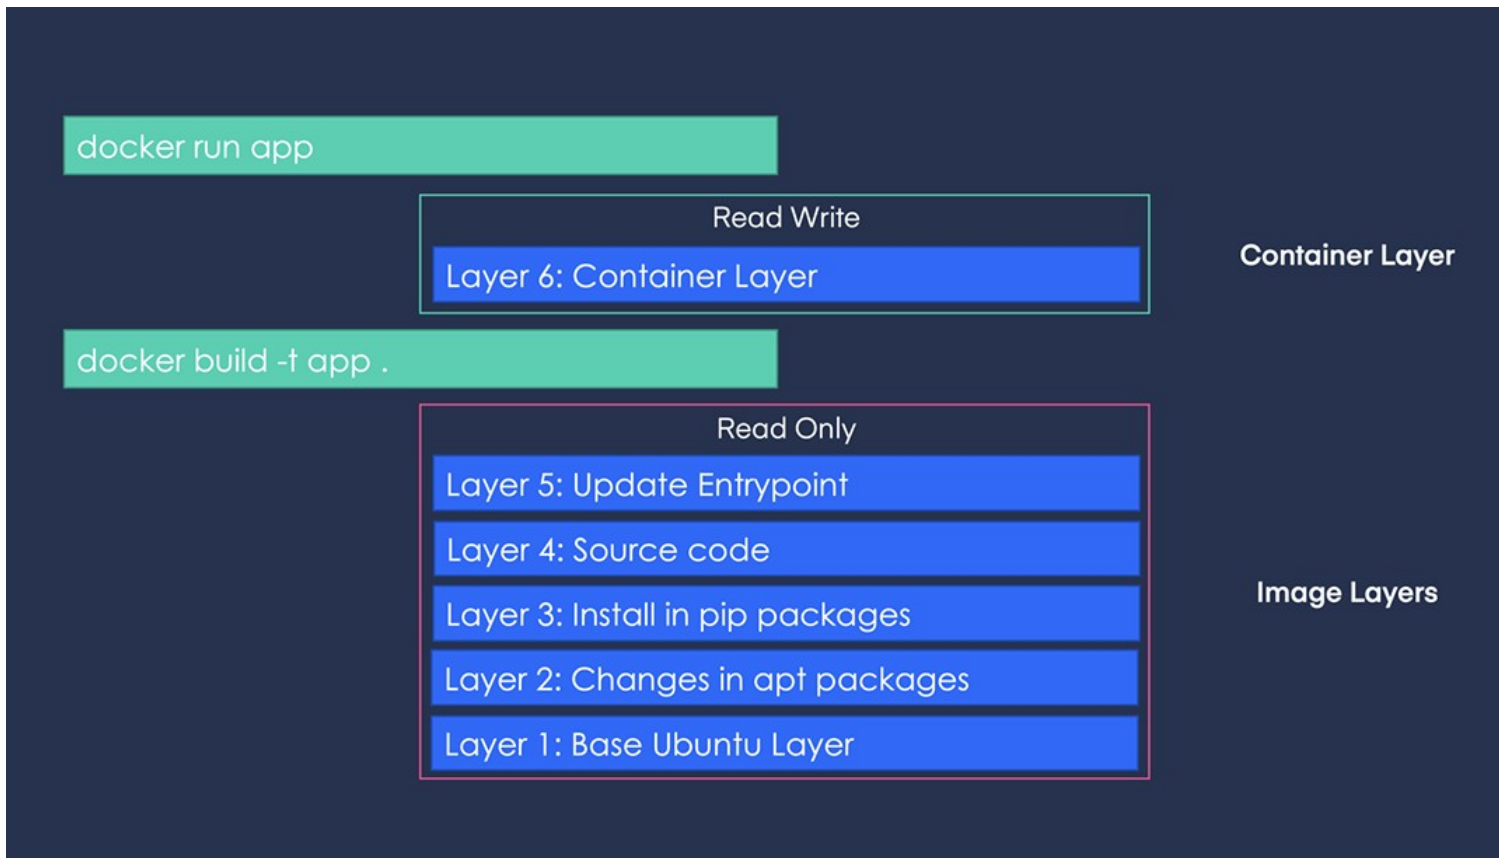
# Docker Images

## Lab Docker 1



# Docker Containers

- Containers are running copies of a Docker image
- Containers have an additional write-able layer added to the image layers



# Running Containers

- The **docker run** command starts a container based on an image
- The image contains a default command to run when it starts
  - Once the command completes, the container exits
  - Some containers, running a web server for example, do not exit
  - These have to be shut down with the **docker stop** command or **docker kill**
- Stopped containers are not destroyed but can be restarted
  - The command **docker start** restarts a stopped container but not exited containers
  - The command **docker create** creates a container but does not run it
  - The command **docker run** = **docker create** followed by **docker start**
- Specific commands inside a container can be executed
  - Inside an already running container with **docker exec <cmd>**
  - Or by starting up a container with **docker run <cmd>**



# Running Containers

- Running containers have a hash id just like images
  - They also can have an optional name **docker run ubuntu --name zippy**
  - Containers are assigned default names otherwise
- There can be multiple containers created from a single image
- Commands used to work with containers
  - **docker ps** – lists all the running processes related to containers
  - **docker ps -a** – lists all of the running and exited processes related to containers
  - Using **docker container ls (-a)** gives exactly the same output
- Docker keeps logs of all activity in each container
  - We can access both a container's logs and monitor its running processes

# Running Containers

- Interactive terminal connections allow us to work within a container if the container supports a shell
  - To work with a shell in Ubuntu we could run **docker run -it ubuntu**
    - If we omit the **-it**, the shell will start up and immediately exit
  - We can override the default command in the image
    - Normally the nginx image starts a web server and does not exit
    - We can start a shell instead with the following command **docker run -it nginx bin/sh**
  - For a running container, we can execute a command using **docker exec -it <container id>**



# Running Docker Containers

Demo





# Running Docker Containers

## Docker Lab 2





# Docker Networks

- Docker engine runs a set of private networks
  - Each container gets an IP address on the docker network
  - If the container provides a service, it normally is exposed through a port on the container
  - The docker engine will map ports on the host networks to ports on containers
- Private docker networks
  - Allow containers to run without interfering with IP addresses or ports on the host system
  - Network types:
    - **Bridge**: the default - creates a private internal isolated network for containers
    - **Host**: allows containers to run on the host network – only implemented for Linux hosts
    - **Overlay**: allows containers on different hosts to communicate with each other
- Exposing ports
  - Services offered by container are specified by port numbers which are made available via port exposing
  - Specific ports to be expose can be defined in an image or a container

# Docker Networking

## Demo





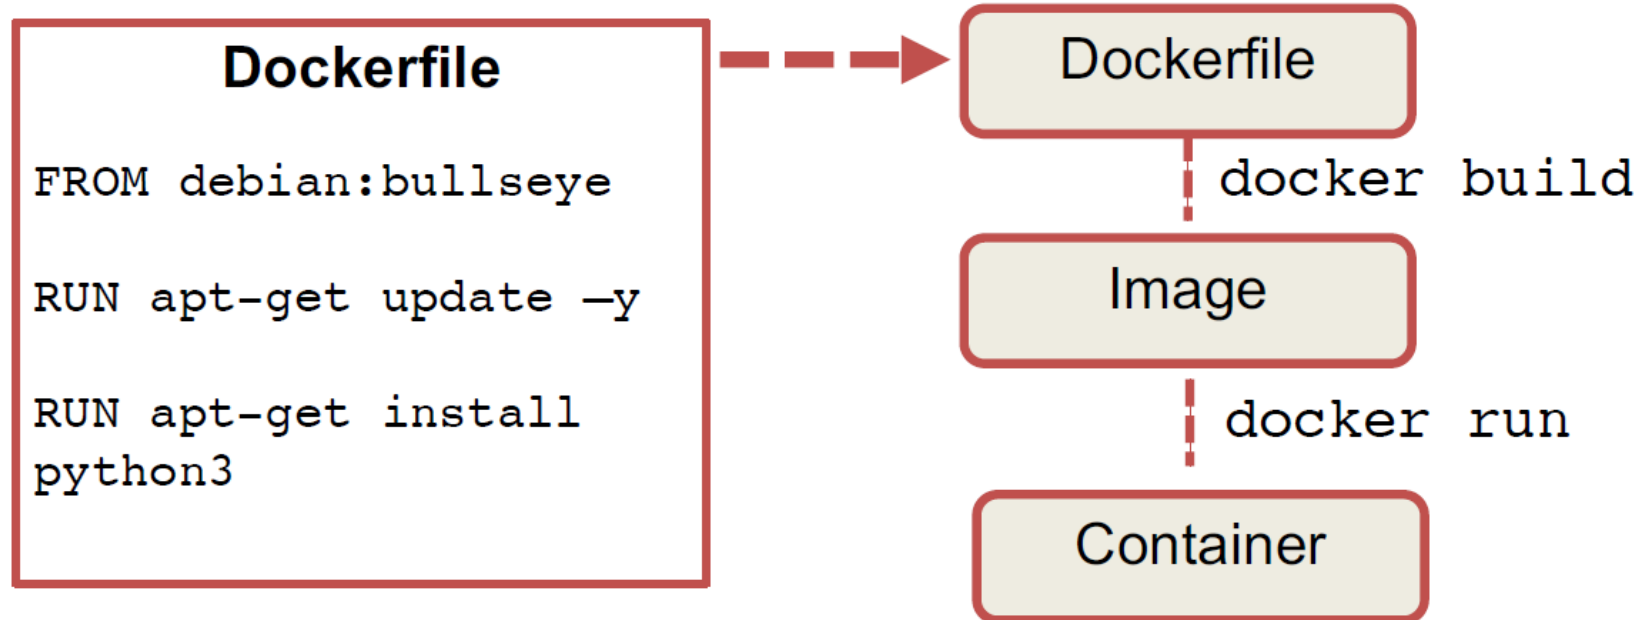
# Docker Networking

## Docker Lab 3



# Building Images with Dockerfile

- Docker can build images automatically by executing instructions in a Dockerfile
  - For example, to build an image with Python three installed on a Debian Linux container





# Dockerfile


- The Dockerfile is a text file
  - contains the instructions that you would execute on the command line to create an image
  - Docker provides a set of standard instructions to be used in the Dockerfile

Command	Description
#	Comment line
MAINTAINER	Provides name and contact info of image creator
FROM	Tells Docker which base image to build on top of (e.g. centos7)
COPY	Copies a file or directory from the build host into the build container
RUN	Runs a shell command inside the build container
CMD	Provides a default command for the container to run. May be overridden or changed
ADD	Copies new files, directories or remote file URLs
LABEL	Adds metadata to an image

# The Build Process

- Docker builds an image by running a series of containers
  - The FROM base image becomes the first layer in the new image
  - The base image is run in a temporary container
  - The first directive, RUN for example, executes and results written to the container's writeable layer
  - The container is committed to a new temporary image
  - The writeable layer now a new image layer
  - This new temporary image is run in a new container, the next directive executes
  - The container is committed to a new temporary image
  - And this continues until the whole Dockerfile is executed

# The Build Process





# Dockerfile

## Demo





# Dockerfiles

## Docker Lab 4





# Docker Monitoring Tools

- Debugging containers can be problematic
- Docker has a number of monitoring tools that can be used
  - `docker logs <container>`: displays console output of the container
  - `docker top <container>`: lists all the processes running in a container
  - `docker stats <container>`: streams real time stats of containers
  - `docker inspect <container>`: displays detailed container configuration

# Docker Monitoring Tools

Demo





# Docker Volumes

- Volumes are where a directory on the host file system is mounted inside a container
  - Anything written to the volume remains on the host file system
  - Any other container can mount the volume and read what was written
- A wide range of file systems can be mounted
  - Amazon AWS buckets for example
- On Windows, there is not direct access to the underlying file since it is created in the Linux VM
  - However, deleting the volume will delete the underlying file



# Docker Volumes

## Demo





# Docker Volumes

## Docker Lab 5





# Docker Compose

- Docker is written in GO
- Docker compose is a developer tool written in Python
- Docker compose is NOT intended to be a production deployment tool
  - It is designed to allow developers to quickly deploy an application using multiple docker containers with single command
  - And to shut it down with a single command
- The configuration of the application is in a docker-compose.yaml file
- Docker compose creates a new private network to run the app in
  - It shuts down the network when the app is shut down
  - Allows multiple instance of the application to run independently and concurrently

# Docker Compose File

- The docker compose file specifies:
  - A set of images to be run called services
  - Any volumes that need to be created or previous volumes to be remounted
  - The configuration information such as ports exposed, etc.
  - Dependencies between the services
- Multiple copies of the application can run at the same time
  - Each instance runs in its own network address space



# Docker Compose File

- The docker compose file specifies:
  - A set of images to be run called services
  - Any volumes that need to be created or previous volumes to be remounted
  - The configuration information such as ports exposed, etc.
  - Dependencies between the services
- Multiple copies of the application can run at the same time
  - Each instance runs in its own network address space

# Docker Compose

## Demo





# Docker Compose

## Docker Lab 6





- **Registries**
  - Tags are of the form `registry_host:port/repository_name:version`
  - Registry services expose the given port in order to manage image requests
  - We will use the registry docker image to provide a registry service on our local machine
  - The docker registry container does not store images, it processes pull and push requests
- **The actual storage location of the images is normally specified**
  - We are using a default of the local file system
  - But it can be a variety of storage mediums, like AWS S3 buckets for example
- **The host running the registry container is referred to as the registry host**
  - We can have any number of registries running, either on different hosts, or the same host but at different ports
  - In our lab, localhost will be our registry server which will be listening on port 5000
  - The registry container mounts the actual physical storage of the images as a docker volume

## Docker Registries

- There are two default registries
- The local cache
  - This is not an actual registry, but where images pulled from different registries are cached
  - Also where our newly created images are stored
  - After creating a new image, it is a best practice to push it to a registry
- When we execute a command like “pull ubuntu”
  - Since there is no registry specified, Docker first checks the local cache
  - If there is no image in the local cache, Docker rewrites the pull request as
    - **Docker pull library/alpine:latest** where “library” is the default repo, in this case docker-hub
  - If there is a registry component to the image tag, that is used instead of the default



- We can download and run the image for the registry executable with
  - **docker run -d -p 5000:5000 --restart=always --name registry registry:2**
- To push an image, we have to tag it with the appropriate registry info
  - **docker tag ubuntu:latest localhost:5000/my-ubuntu**
  - Notice our registry is the host and port, the name “registry” is just the docker container that processes the image requests into and out of the associated docker volume
- The we just push it to the registry
  - **docker push localhost:5000/my-ubuntu**
  - Now there exist two copies of the image, one in the cache and one in the registry
  - Normally we use a registry name that is self explanatory like:
    - **docker run -d -p 5000:5000 --restart=always --name devteam1 registry:2**
    - **docker run -d -p 5500:5000 --restart=always --name prod\_qa registry:2**
  - But docker commands use the port and host, so docker sees these as localhost:5000 and localhost:5500

- Normally there is no GUI to a private registry
  - However there are web based tools that can be used
- The standard commands work on the cache
  - If you can't execute a command line command on docker hub, you can't do it on your registry
- The local cache is where we do our docker programming
- In brief
  - A registry is a docker volume
  - The registry container is a service that runs in a container that pulls and pushes images into and out of the volume
  - When we remove a registry, it is not enough to remove the registry container
  - We also have to remove the associated volume



# Docker Registries

## Demo





# Docker Registries

## Docker Lab 7





A detailed Renaissance-style painting of Plato's Academy. The scene is set in a grand hall with classical columns and a landscape view through the arches. Numerous figures, representing various Greek philosophers, are engaged in discussion and study. Plato, an older man with a white beard, stands in the center, pointing his right index finger towards the sky. Aristotle, a younger man with a dark beard, stands next to him, gesturing with his palm facing down towards the earth. Other figures are seated or standing, some holding books or scrolls, and others gesturing in conversation. The composition is dense and dynamic, with a strong sense of perspective.

**Questions?**

# Class Project Discussion







End Module